



Anomaly Detection in High Dimensional Data

Priyanga Dilini Talagala, Rob J. Hyndman & Kate Smith-Miles

To cite this article: Priyanga Dilini Talagala, Rob J. Hyndman & Kate Smith-Miles (2020): Anomaly Detection in High Dimensional Data, Journal of Computational and Graphical Statistics, DOI: [10.1080/10618600.2020.1807997](https://doi.org/10.1080/10618600.2020.1807997)

To link to this article: <https://doi.org/10.1080/10618600.2020.1807997>



View supplementary material [↗](#)



Accepted author version posted online: 13 Aug 2020.



Submit your article to this journal [↗](#)



Article views: 10



View related articles [↗](#)



View Crossmark data [↗](#)

Anomaly Detection in High Dimensional Data

Priyanga Dilini Talagala^{1,3,4}

and

Rob J. Hyndman^{1,3}

and

Kate Smith-Miles^{2,3}

¹Department of Econometrics and Business Statistics, Monash University, Australia

²School of Mathematics and Statistics, University of Melbourne, Australia

³ARC Centre of Excellence for Mathematics and Statistical Frontiers (ACEMS), Australia

⁴Department of Computational Mathematics, University of Moratuwa, Sri Lanka

Corresponding author Priyanga Dilini Talagala priyangad@uom.lk

Abstract

The HDoutliers algorithm is a powerful unsupervised algorithm for detecting anomalies in high-dimensional data, with a strong theoretical foundation. However, it suffers from some limitations that significantly hinder its performance level, under certain circumstances. In this article, we propose an algorithm that addresses these limitations. We define an anomaly as an observation where its k -nearest neighbour distance with the maximum gap is significantly different from what we would expect if the distribution of k -nearest neighbours with the maximum gap is in the maximum domain of attraction of the Gumbel distribution. An approach based on extreme value theory is used for the anomalous threshold calculation. Using various synthetic and real datasets, we demonstrate the wide applicability and usefulness of our algorithm, which we call the stray algorithm. We also demonstrate how this algorithm can assist in detecting anomalies present in other data structures using feature engineering. We show the situations where the stray algorithm outperforms the HDoutliers algorithm both in accuracy and computational time. This framework is implemented in the open source R package `stray`.

Keywords: Extreme value theory, High dimensional data, Nearest neighbour searching, Temporal data, Unsupervised outlier detection

1 Introduction

The problem of anomaly detection has many different facets, and detection techniques can be highly influenced by the way we define anomalies, type of input data and expected output. These differences lead to wide variations in problem formulations, which need to be addressed through different analytical techniques. Although several useful computational methods currently exist, developing new methods for anomaly detection continues to be an active, attractive interdisciplinary research area owing to different analytical challenges in various application fields. Ever-increasing computing resources and advanced data collection technologies that emphasise real-time, large-scale data are other reasons for this growth since they introduce new analytical challenges with their increasing size, speed and complexity that demand effective, efficient analytical and computing techniques.

Anomaly detection has two main objectives, which are conflicting in nature: One downgrades the value of anomalies and attempts eliminating them, while the other demands special attention be paid to anomalies and root-cause analysis be conducted. The presence of anomalies in data can be considered data flaws or measurement errors that can lead to biased parameter estimation, model misspecification and misleading results if classical analysis techniques are blindly applied (Ben-Gal 2005, Abuzaid et al. 2013). In such situations, the focus is to find opportunities to remove anomalous points and thereby improve both the quality of the data and results from the subsequent data analysis (Novotny and Hauser 2006). In contrast, in many other applications, anomalies themselves are the main carriers of significant and often critical information, such as extreme environmental conditions (e.g., bushfire, tsunami, flood, earthquake, volcanic eruption and water contamination), faults and malfunctions (e.g., flight tracking

and power cable tracking) and fraud activities ([Ben-Gal 2005](#)), that can cause significant harm to valuable lives and assets if not detected and treated quickly.

High-dimensional datasets exist across numerous fields of study ([Liu et al. 2016](#)). Some anomaly detection algorithms also use feature engineering as a dimension reduction technique and thereby convert other data structures, such as a collection of time series using time series features ([Talagala, Hyndman, Leigh, Mengersen and Smith-Miles 2019](#), [Hyndman et al. 2015](#)), collection of scatterplots using scagnostics ([Wilkinson et al. 2005](#)) and genomic micro arrays and chemical compositions in biology ([Liu et al. 2016](#)) into high-dimensional data prior to the detection process for easy control. Under the high-dimensional data scenario, all attributes can be of the same data type or a mixture of different data types, such as categorical or numerical, which has a direct impact on the implementation and scope of the algorithm. Much research attention has been paid to anomaly detection for numerical data ([Breunig et al. 2000](#), [Tang et al. 2002](#), [Jin et al. 2006](#), [Gao et al. 2011](#)). Limited methods are available that treat both numerical and categorical data using correspondence analysis, for example, as in [Wilkinson \(2017\)](#).

High-dimensional anomalies can arise in all the attributes or a subset of the attributes ([Unwin 2019](#)). If all anomalies in a high-dimensional data space were anomalies in a lower dimension, then anomaly detection can be performed using axis parallel views or by incorporating an additional step of variable selection for the detection process. However, in practice, certain high-dimensional instances are only perceptible as anomalies if treated as high-dimensional problems and the correlation structure of all the attributes considered. Otherwise, these tend to be overlooked if attributes are considered separately ([Wilkinson 2017](#), [Ben-Gal 2005](#)).

The problem of anomaly detection has been extensively studied over the past decades in many application domains ([Chandola](#)

et al. 2009, Aggarwal 2017, Shahid et al. 2015, Gupta et al. 2014, Hodge and Austin 2004). Among the many possibilities, the HDoutliers algorithm, recently proposed by Wilkinson (2017), is a powerful unsupervised algorithm, with a strong theoretical foundation, for detecting anomalies in high-dimensional data. The study presented by Talagala, Hyndman, Leigh, Mengersen and Smith-Miles (2019) also verifies its performances through a thorough comparative evaluation of existing state-of-the-art anomaly detection methods. Although this algorithm has many advantages, a few characteristics hinder its performance. In particular, under certain circumstances it tends to increase the rate of false negatives (i.e., the detector ignores points that appear to be real anomalies) because it uses only the nearest-neighbour distances to distinguish anomalies. Further, to deal with large datasets with numerous observations it uses the Leader algorithm (Hartigan and Hartigan 1975), which forms several clusters of points in one pass through the dataset using a ball of a fixed radius. By incorporating this clustering method, it tries to gain the ability to identify anomalous clusters of points. However, in the presence of very close neighbouring anomalous clusters it tends to increase the rate of false negatives. Further, this additional step of clustering has a serious negative impact on the computational efficiency of the algorithm when dealing with large datasets.

Through this study, we make three fundamental contributions. First, we propose an algorithm called *stray*, representing ‘**Search and TRace AnomaLY**’, that addresses the limitations of the HDoutliers algorithm. The *stray* algorithm presented here focuses specifically on fast, accurate anomalous score calculation using simple but effective techniques for improved performance. Second, we introduce an R (R Core Team 2019) package, *stray* (Talagala, Hyndman and Smith-Miles 2019), that implements the *stray* algorithm and related functions. Third, we demonstrate the wide applicability and usefulness of our *stray* algorithm, using various datasets.

Our improved algorithm, stray, has many advantages: (1) It can be applied to both one-dimensional and high-dimensional data. (2) It is unsupervised in nature and therefore does not require training datasets for the model-building process. (3) The anomalous threshold is a data-driven threshold and has a valid probabilistic interpretation because it is based on the extreme value theory. (4) The presence of one anomalous instance can “mask” the appearance of another anomalous instance (Hadi 1992). The stray algorithm deals with the masking problem by using k -nearest neighbour distances for anomalous score calculations. (5) It can provide near real-time support to datasets that stream in large quantities owing to its use of fast nearest neighbour searching mechanisms. (6) It can deal with data that may have multimodal distributions for typical data instances. (7) It produces both score (to indicate how anomalous the instances are) and binary classification (to reduce the searching space during the visual and root-cause analysis) for each data instance as an output. (8) It can detect outliers as well as inliers.

The stray algorithm differs substantially from the HDoutliers algorithm in the following ways. (1) It does not determine the outliers solely with the help of the nearest neighbour distances; instead, the algorithm considers the k -nearest neighbour distance with the maximum gap for all observations and defines an anomalous threshold on the k nearest neighbour distances with the maximum gap. (2) The HDoutliers algorithm deals with the masking problem by incorporating a clustering step which is extremely time-consuming, particularly with large samples in high dimensions. Owing to the use of k -nearest neighbour distances, the stray algorithm does not require an additional computationally-intensive clustering procedure to detect anomalous clusters of points. (3) The HDoutliers algorithm provides only a binary classification; in contrast, the stray algorithm produces both a binary classification and an anomalous score indicating the degree of outlierness of each measurement. (4) The simulation study shows that the stray implementation is much faster than the HDoutliers package implementation, particularly with large samples in high dimensions.

The remainder of this paper is organised as follows. Section 2 presents the related work to lay the foundation for the stray algorithm. Section 3 describes the limitations of the HDoutliers algorithm that hinder its performance. Section 4 presents the improved algorithm, stray, that addresses the limitations of the HDoutliers algorithm. Section 5 presents a comprehensive evaluation, illustrating the key features of the stray algorithm. Section 6 includes an application of stray algorithm related to pedestrian behaviour in the city of Melbourne, Australia. Section 7 concludes the article and presents future research directions.

2 Background

2.1 Types of Anomalies in High Dimensional Data

The problems of anomaly detection in high-dimensional data are threefold, involving detection of: (a) global anomalies, (b) local anomalies and (c) micro clusters or clusters of anomalies ([Goldstein and Uchida 2016](#)). Most of the existing anomaly detection methods for high-dimensional data can easily recognise global anomalies since they are very different from the dense area with respect to their attributes. In contrast, a local anomaly is only an anomaly when it is distinct from, and compared with, its local neighbourhood. [Madsen \(2018\)](#) introduces a set of algorithms based on a density or distance definition of an anomaly, which mainly focuses on local anomalies in high-dimensional data. Micro clusters or clusters of anomalies may cause masking problems. Very little attention has been paid to this problem relative to the other two categories. The recently proposed HDoutliers algorithm ([Wilkinson 2017](#)) addresses this problem to some extent by grouping instances together that are very close in the high-dimensional space and then selecting a representative member from each cluster before calculating nearest neighbour distances for the selected instances. In this study, we focus on all three of these anomaly types.

2.2 Definitions for Anomalies in High Dimensional Data

Owing to the complex nature of the problem, it is difficult to find a unified definition for an anomaly and the definition often depends on the focus of the study and the structure of the input data available to the system (Williams 2016, Unwin 2019). However, there are some definitions that are general enough to cope with datasets with various application domains. Grubbs (1969) defines an anomaly as an observation that deviates markedly from other members of the dataset. However, this deviation can be defined in terms of either distance or density. Burridge and Taylor (2006), Wilkinson (2017) and Schwarz (2008) have all proposed methods for anomaly detection by defining an anomaly in terms of distance. In contrast, Hyndman (1996), Clifton et al. (2011) and Talagala et al. (2020) have proposed methods that define an anomaly with respect to either the density or the chance of the occurrence of observations. Madsen (2018) also provides a series of distance and density-based anomaly detection algorithms.

In this study, we define an anomaly as an observation where its k -nearest neighbour distance with the maximum gap is significantly different from what we would expect if the corresponding distribution is in the maximum domain of attraction of the Gumbel distribution. This covers a wide range of distributions. Intuitively, an anomaly deviates markedly from the majority of the data, with a significantly larger distance between typical observations and anomalies compared to the distances between typical observations. This definition allows anomalies to be isolated observations or small isolated clusters of observations.

3 Limitations of HDoutliers Algorithm

The HDoutliers algorithm (Wilkinson 2017) is a distance based anomaly detection algorithm. One important property of this algorithm is that it has an ability to convert any higher dimensional anomaly detection problem to a one dimensional problem by taking the nearest neighbour distances of the data instances.

There are two published versions of the HDoutliers algorithm. The first version calculates nearest neighbour distance for each data instance and does not involve any clustering step prior to the nearest neighbour distance calculation. This version of the algorithm (version 1 of the HDoutliers, hereafter) is recommended for small samples. The default maximum sample size for version 1 is set to 10000 in the R implementation of the `HDoutliers` package. The second version of the HDoutliers algorithm incorporates a clustering step with the aim of detecting micro clusters. It uses the Leader algorithm ([Hartigan and Hartigan 1975](#)) to form several clusters of points and then selects a representative member from each cluster. The nearest neighbour distances are then calculated only for the selected representative members.

Although the HDoutliers algorithm ([Wilkinson 2017](#)) has many advantages, a few characteristics limit its possibilities. Next, we discuss these limitations in detail.

3.1 HDoutliers Uses Only the Nearest Neighbour Distance to Discriminate Anomalies

The HDoutliers algorithm uses the Leader algorithm to form small clusters of points, prior to calculating nearest neighbour distance. In the Leader algorithm, each cluster is a ball in the high-dimensional data space. In the HDoutliers algorithm, the radius of this ball is selected such that it is well below the expected value of the distances between $n(n-1)/2$ pairs of points distributed randomly in a d -dimensional unit hypercube.

After forming clusters using the Leader algorithm, the HDoutliers algorithm selects representative members from each cluster. It then calculates the nearest neighbour distances for each of these representative members. These distances are then used to identify the anomalies based on the assumption that anomalies bring large distance separations between typical data and the anomalies, in comparison to the separations between typical data themselves. Therefore, under this assumption it is believed that any anomalous cluster will appear far

away from the clusters of the typical data points. As a result, the nearest neighbour distance for this anomalous cluster will be significantly higher than that of the clusters of typical data and thereby identify it as an anomalous cluster. All the data points contained in the anomalous cluster are then marked as anomalous points within a given dataset.

However, one further assumption for this method to work properly is that any anomalous clusters present in the dataset are isolated. For example, imagine a situation in which two anomalous clusters are very close to one another but are far away from the rest of the typical clusters. Now, the two clusters will become nearest neighbours to one another and they will jointly protect them by being anomalous by giving very small nearest neighbour distances for both clusters that are compatible with the nearest neighbour distances of the rest of the typical clusters. Figures 4 (c-II) and (d-II) further elaborate this argument. In these two examples, the HDoutliers algorithm (with the clustering step) declares points as anomalies only if they are isolated and fails to detect anomalous clusters that share a few cluster neighbours. Although the HDoutliers algorithm incorporates the clustering step with the aim of identifying anomalous clusters of points, because of the very small size of the ball that is used to produce clusters (exemplars) in the d -dimensional space, it fails to bring all the points into a single cluster and instead produces a few anomalous clusters that are very close to one another. These anomalous clusters then become nearest neighbours to one another and have very small nearest neighbour distances for the representative member of each cluster. Since the detection of anomalies entirely depends on these nearest neighbour distances and since the anomalous clusters do not show any significant deviation from typical clusters with respect to the nearest neighbour distances, the algorithm now fails to detect these points as anomalies and thereby increases the rate of false negatives.

3.2 Problems Due to Clustering Via Leader Algorithm

After forming clusters of data points, the HDoutliers algorithm completely ignores the density of the data points. Once it forms clusters of data points using the Leader algorithm, it selects a representative member from each cluster and carries out further analysis only using these representative members. Figure 4 (e-II) provides an example related to this issue. This dataset is a bimodal dataset with an anomalous point located between the two typical classes. The entire dataset contains 2,001 data points. The data points gathered at the leftmost upper corner represent one typical class with 1,000 data points. The second typical class of data points is gathered at the rightmost bottom corner with another 1,000 data points. Since this second class of data points is closely compacted in substance, the 1,000 data points are now wrapped by a single ball when forming clusters using the Leader algorithm. In the HDoutliers algorithm, the next step is to select one member from each of these clusters. Once it selects a representative member from this ball that contains 1,000 data points, it ignores the remaining 999 data points in detecting anomalies. This step misleads the algorithm, and the remaining steps of the algorithm view this representative member as an isolated data point, although it is surrounded by 999 neighbouring data points in the original dataset. Therefore, all data points in this entire class are declared as anomalies by the algorithm, although it contains half of the dataset. Unwin (2019) suggests jittering not as a perfect solution, but as an alternative to mitigate this problem. Unwin (2019) also argues that the problem tends not to occur in high-dimensional data spaces where this kind of granularity is less likely. However, then it gives rise to the problem of neighbouring anomalous clusters (as illustrated in Figure 4 (c-II, d-II)), which individually appear to be typical, or of limited suspicion (due to the presence of other neighbouring anomalous clusters), yet, their co-occurrence is highly anomalous.

Figure 4 (f-II) provides another situation in which false negatives increase because of the clustering step. This bivariate dataset contains 1,001 data points. The data points gathered at the leftmost upper corner represent a typical class covering 1,000 data points, and the isolated data point at the rightmost bottom

corner represents an anomaly. Since this typical class of 1,000 data points is closely compacted, it gives rise to only 14 clusters through the Leader algorithm. Altogether, the dataset forms 15 clusters with the one created by the isolated point located at the rightmost bottom corner. Even though the original dataset contains 1,001 data points, the algorithm considers only 15 data points (a representative member from each cluster) for calculating the anomalous threshold. Now, this number is not large enough to yield a stable estimate for the anomalous threshold. Due to this ignorance of the density of the original dataset, it now fails to detect the obvious anomalous point at the leftmost bottom corner.

3.3 Problem with Threshold Calculation

A companion R package ([Fraley 2018](#)) is available for the algorithm proposed by [Wilkinson \(2017\)](#). According to the R package implementation, the current version of the HDoutliers algorithm uses the next potential candidate for anomalies in calculating the anomalous threshold, in each iteration of the bottom-up searching algorithm. This approach causes an increase in the false detection rate under certain circumstances. We avoid this limitation in our proposed algorithm.

4 Proposed Improved Algorithm: stray Algorithm

In this section, we propose an improved algorithm for anomaly detection in high dimensional data. Our proposed algorithm is intended to overcome the limitations of the HDoutliers algorithm and thereby enhance its capabilities.

4.1 Input to the stray Algorithm

An input to the stray algorithm is a collection of data instances where each data instance can be a realisation of only one attribute or a set of attributes (also referred to using terms such as features, measurements and dimensions). In this study, we limit our discussion to quantitative data; therefore, an input can be a vector, matrix or data frame of $d(\geq 1)$ numerical variables, where each column

corresponds to an attribute and each row corresponds to an observation of these attributes. The focus is then to detect anomalous instances (rows) in the dataset. The stray algorithm can be easily extended to deal with categorical data by using correspondence analysis which converts categorical data into quantitative information, as in [Wilkinson \(2017\)](#).

4.2 Normalise the Columns

Since the stray algorithm is based on the distance definition of an anomaly, a distance measure must be specified. In this analysis we use Euclidean distances to measure the k -nearest neighbour distances between observations in the high-dimensional data space. To make the variables of equivalent weight, the columns of the data are first normalised using a linear transformation. By default, we use “min-max normalisation”, with the resulting data ranging from 0 to 1. Alternative normalizations based on linear transformations (see [Kandanaarachchi et al. 2018](#)) may also be used and are available through the `stray` package implementation. Thus, our distance measure is invariant to rescaling the variables but it is not affine invariant (unlike Mahalanobis distance, for example). Exploration of alternative distance measures would be an interesting avenue for future research. In addition to min-max normalisation, a robust normalisation method ($(x - \text{median}(x)) / \text{IQR}(x)$) is also available through the `stray` package implementation. However, there is no one-fit-for-all normalisation strategy for anomaly detection problems even though min-max normalisation is shown to be preferred to median-IQR with most of the datasets and anomaly detection methods considered in [Kandanaarachchi et al. \(2018\)](#).

4.3 Nearest Neighbour Searching

In the stray algorithm, after the columns of the dataset are normalised, it calculates the Euclidean distance-based k -nearest neighbour distance with the maximum gap for each and every instance. By using this measure, we were able to address the aforementioned limitations of the HDoutliers algorithm.

For each individual observation, the algorithm first calculates the k -nearest neighbour distances, $d_{i,KNN}$, where $i = 1, 2, \dots, k$. Then, it calculates the successive differences between distances, $\Delta_{i,KNN}$. Next, it selects the k -nearest neighbour distance with the maximum gap, $\Delta_{i,max}$. Figure 1 illustrates how these steps help our improved algorithm to detect anomalous points or anomalous clusters of points.

In Figure 1 (a), the dataset contains only one anomaly at (15,16.5). For this dataset, the nearest neighbour distance can differentiate the anomalous point from the remaining typical points because the nearest neighbour distance for the anomalous point is significantly larger (14.8) than that for the remaining typical points. Figure 1 (b) shows the change in the k -nearest neighbour distances of the anomaly at (15,16.5). For this dataset, the k -nearest neighbour distance with the maximum gap occurs when $k = 1$. The second dataset, in Figure 1 b), has a micro cluster with three anomalies around (15,16.5). If only the nearest neighbour distances are calculated for each observation, then the three anomalies in the micro cluster are not distinguishable from the typical points since their values are very small (0.7) compared with that of most typical points with nearest neighbour distances at around (0.0015 to 2.5). However, the three anomalies in the micro cluster are distinguishable from their typical points with respect to the k -nearest neighbour distances with the maximum gap at around 14.8 and 14.9 (Figure 1 (d)). For the three anomalies in the micro cluster in Figure 1 (d), the third nearest neighbour distance has the maximum gap (Figure 1 e)) and the three points in the micro cluster are now easily distinguished as anomalies, with respect to k -nearest neighbour distances with the maximum gap. Therefore, by using k -nearest neighbour distances with the maximum gap, the stray algorithm gains the ability to detect both anomalous singletons and micro clusters. Through this approach, we are able to reduce the false detection rate and thereby address the limitations of the HDoutliers algorithm, while gaining the ability to detect micro clusters. This is also a very simple, but clever, investment as compared with the time taken by the leader algorithm to form small clusters to

detect micro clusters (especially for datasets with large dimensions), in the HDoutliers algorithm. Further, for each point, the corresponding k -nearest neighbour distances with the maximum gap act as an anomalous score to indicate the degree of being an anomaly.

In the current study, we consider both exact and approximate k -nearest neighbour searching techniques. Brute force search involves going through every possible pairing of points to detect k -nearest neighbours for each data instance, and therefore, exact k -nearest neighbours are explored. Conversely, k -dimensional trees (k-d trees) employ spatial data structures that partition space to allow efficient access to a specified query point (Elseberg et al. 2012 a). Therefore, it involves searching approximate k -nearest neighbours around a specified query point.

In the current algorithm, parameter k , which determines the size of the neighbourhood, is introduced as a user-defined parameter that can be selected according to the application. One way to interpret the role of k in the stray algorithm is to view it as the minimum possible size for a typical cluster in a given dataset. If the size of an anomalous cluster is less than k , it will be detected as a micro cluster by the stray algorithm. The choice of k has different effects across different dimensions and sizes of data (Campos et al. 2016). We can set k to 1 if no micro clusters are present in the dataset and thereby focus on local and global anomalous points. High k values are recommended for datasets with high dimensions because of the curse of dimensionality.

4.4 Threshold Calculation

Anomalous scores assign each point a degree of being an anomaly. However, for certain applications it is also important to categorise typical and anomalous points for the subsequent root-cause analysis. Ideally, we prefer a universal threshold to unambiguously distinguish anomalous points from typical points. Following Schwarz (2008), the HDoutliers algorithm (Wilkinson 2017) defines an

anomalous threshold based on extreme value theory, a branch of probability theory that relates to the behaviour of extreme order statistics in a given sample (Galambos et al. 2013).

The anomalous threshold calculation in Schwarz (2008); Burridge and Taylor (2006) and Wilkinson (2017) is an application of Weissman's spacing theorem (Weissman 1978) (Theorem 1) that is applicable to the distribution of data covered by the maximum domain of attraction of a Gumbel distribution. This requirement is satisfied by a wide range of distributions, ranging from those with light tails to moderately heavy tails that decrease to zero faster than any power function (Embrechts et al. 2013). Examples include the exponential, gamma, normal and log-normal distributions with exponentially decaying tails.

Let X_1, X_2, \dots, X_n be a sample from a distribution function F and let $X_{1:n} \geq X_{2:n} \geq \dots \geq X_{n:n}$ be the order statistics. The available data are $X_{1:n}, \dots, X_{K:n}$ for some fixed K .

Theorem 1 (Spacing Theorem). (*Proposition 1 in Burridge and Taylor (2006), p.6 and Theorem 3 in Weissman (1978), p.813; the notations have been changed for consistency in this paper*)

Let $D_{i,n} = X_{i,n} - X_{i+1,n}$, ($i = 1, \dots, K$) be the spacing between successive order statistics. If F is in the maximum domain of attraction of the Gumbel distribution, the spacings $D_{i,n}$ are asymptotically independent and exponentially distributed with mean proportional to i^{-1} .

We illustrate this theorem using Figure 2, which shows the distribution of the descending order statistics ($X_{i,n}$) and the standardised spacings, ($iD_{i,n}$), for $i \in \{1, \dots, 10\}$ for 1, 000 samples each containing 20, 000 random numbers from the standard normal distribution. Figure 2 (a) shows the distribution of $X_{i,n}$ with means of $X_{i,n}$ depicted as black crosses. The gaps between consecutive black crosses give the spacings between higher-order statistics ($D_{i,n}$). We note that

the normal distribution is in the maximum domain of attraction of the Gumbel distribution and that this example contains no outliers. A consequence of Theorem 1 is that the standardised spacings $(iD_{i,n})$ for $(i = 1, \dots, K)$, are approximately iid (Burridge and Taylor 2006). Figure 2 (b) shows the distribution of the standardised spacings $(iD_{i,n})$ for $(i = 1, 2, \dots, 10)$ for 1, 000 samples of size 20, 000. Each letter-value box plot (Hofmann et al. 2017) exhibits approximately the shape of an exponential distribution.

Following Schwarz (2008), Burridge and Taylor (2006) and Wilkinson (2017), we start our anomalous threshold calculation from a subset of the points covering 50 per cent of them with the smallest anomalous scores (k nearest neighbour distances with the maximum gap) under the assumption that this subset contains the anomalous scores corresponding to typical data points and the remaining subset contains the scores corresponding to the possible candidates for anomalies. Following the Weissman spacing theorem, it then fits an exponential distribution to the upper tail of the outlier scores of the first subset, and then computes the upper $1 - \alpha$ points of the fitted cumulative distribution function, thereby defining an anomalous threshold for the next anomalous score as follows:

Let X be the k -nearest neighbour distances with the maximum gap for the first subset of the dataset with cdf F and $F \in MDA(\Lambda^+)$. Let $X_{1,n} \geq X_{2,n} \geq \dots \geq X_{n,n}$ be the order statistics. Consider $D_{i,n} = X_{i,n} - X_{i+1,n}$, where $i = 1, \dots, K$ for some fixed K . According to the Weissman spacing theorem, we have

$$D_{i,n} \sim \exp(\lambda i),$$

where $E(D_i) = 1/(\lambda i)$ and $1/\lambda$ is the proportional constant.

Let t be the anomalous threshold. Then $\Pr(D_{1,n} \leq t) = 1 - \alpha$, and so

$$t = \log(1/\alpha) / \lambda.$$

Consider, $\hat{D}_1 = \sum_{i=2}^K (iD_i) / (k-1)$. This is an unbiased estimator for $1/\lambda$. Therefore, an empirical threshold is given by

$$\hat{t} = \log(1/\alpha) \sum_{i=2}^K iD_i / (k-1).$$

From the remaining subset, the point with the smallest anomalous score is selected. If this anomalous score exceeds the empirical threshold \hat{t} , all the points in the remaining subset are flagged as anomalies. Otherwise, it declares the point as a typical point and adds it to the subset of the typical points. It then updates the cut-off point, including the latest addition. This searching algorithm continues until it finds an anomalous score that exceeds the latest cut-off point. This algorithm is known as a 'bottom-up searching' algorithm in [Schwarz \(2008\)](#). This threshold calculation is performed under the assumption that the distribution of k -nearest neighbours with the maximum gap is in the maximum domain of attraction of the Gumbel distribution, which covers a wide range of distributions.

4.5 Output

In stray, anomalies are measured in two scales: (1) binary classification and (2) outlier score. Under binary classification, data instances are classified either as typical or anomalous using the data-driven anomalous threshold based on the extreme value theory. This anomalous threshold or the cut-off point for the anomalies is defined on the k -nearest neighbour distances with the maximum gap. This type of classification is important if the subsequent steps of the data analysis process are automated.

The decision of the stray algorithm is based on the k -nearest neighbour distances with the maximum gap of the data instances. The k -nearest neighbour distance with the maximum gap of each data instance quantifies how isolated the data instance is with respect to its surrounding points, and hence acts as an anomalous score. High values indicate higher degree of outliyness or isolation,

while low values indicate high compactness (Figures 1 (d), (e)). Therefore, in addition to the binary output, the stray algorithm also assigns an anomalous score to each data instance. These anomalous scores allow the user to rank and select the most serious or relevant anomalous points for root-cause analysis and taking immediate precautions. The HDoutliers algorithm (Wilkinson 2017), which provides only a binary classification, does not directly allow the user to make such a choice to direct their attention to more significant anomalous instances. Conversely, various methods proposed in the literature provide anomalous scores, but the anomalous threshold is user defined and application specific (Madsen 2018). The output produced by stray is an all-in-one solution encapsulating necessary measurements of anomalies for further actions.

5 Experiments

The HDoutliers algorithm is a powerful algorithm in the current state-of-the-art methods for detecting anomalies in high-dimensional data. The focus of the stray algorithm is to address some of the limitations of the HDoutliers algorithm that hinder its performance under certain circumstances. Here, we perform an experimental evaluation on the accuracy and computational efficiency of our stray algorithm relative to the HDoutliers algorithm. While these examples are fairly limited in number and are mostly limited to bivariate datasets, they should be viewed only as simple illustrations of the key features of the stray algorithm that outperforms the HDoutliers algorithm.

The first experiment (Figure 3) was designed to test the effect of the dimension, size of the data and the k -nearest neighbour searching method on running times of the different versions of the two algorithms: stray and HDoutliers. Due to high computational cost, the algorithms were applied to problems with only up to 100 dimensions, which are still relatively small for many real-world problems.

Compared with version 1 of the HDoutliers algorithm (Figure 3 (a)), version 2 with the clustering step is extremely slow for higher dimensions (>10), and the running

time increases more rapidly with increasing sample size. For clear comparison between the different versions of the two algorithms (stray and HDoutliers), only a part of the measurements of the full experiment of the second version of the HDoutliers algorithm is displayed in Figure 3 (b-I)). Figure 3 (b-II) presents the full version of Figure 3 (b-I). The additional clustering step in the second version of the HDoutliers algorithm, which is essential for detecting micro clusters, is extremely time-consuming, particularly with large samples with higher dimensions. Figure 3 (c)–(f) corresponds to the stray algorithm. In this experiment, to ascertain the influence from the k -nearest neighbour searching methods, we considered both exact (brute force) and approximate (kd-trees) nearest neighbour searching algorithms.

Many implementations of k -nearest neighbour searching algorithms are available for the R software environment. We considered the `FNN` (Beygelzimer et al. (2019), Figure 3 (c) & (d)) and `nabor` (Elseberg et al. (2012 b); Figure 3 (e) & (f)) R packages for our comparative analysis. R package `nabor`, wraps a fast k -nearest neighbour library written in templated C++. We noticed that searching $k - (> 1)$ nearest neighbours (Figure 3 (a), in this example k is set to 10) instead of only one ($k = 1$) nearest neighbour (Figure 3 (d)) increases the running time only slightly as the number of instances is increased. The results in both Figure 3 (a) and Figure 3 (d) are based on approximate nearest neighbour distances using the kd-trees nearest neighbour searching algorithm. We observed that the kd-trees implementation in the `nabor` package (Figure 3 (f)) is much faster than the `FNN` package implementation (Figure 3 (d)). Surprisingly, as the dimension increases, the running time of the stray algorithm with kd-trees (Figure 3 (d), (f)) increases much more quickly than that of the brute force algorithm, which involves searching every possible pairing of points to detect k -nearest neighbours for each data instance (Figure 3 (c), (e)). Other studies (Kanungo et al. 2002) have also reported a similar result for algorithms based on kd-trees and its variants. This could be due to the parallelisability and memory access patterns of the two searching mechanisms. The brute force algorithm is easily

parallelisable because it involves independent searching of all possible candidates for each data instance. In contrast, the kd-tree searching algorithm is naturally serial and therefore difficult to implement on parallel systems with appreciable speedup (Zhang 2017).

Following Wilkinson (2017), we evaluated the false positive rate (typical points incorrectly identified as anomalies) of the stray algorithm by running it many times on random data. The values presented in Table 1 are based on 1000 iterations and the mean values are reported. Different versions of the two algorithms (stray and Hdoutliers) were applied on datasets where each column is randomly generated from the standardised normal distribution. In each test, the critical value, α , was set to 0.05. Compared with the HDoutliers algorithm, low false positive rates were achieved for the stray algorithm across all dimensions and sample sizes. Unlike in the HDoutliers algorithm (Unwin 2019), in stray a much smaller false detection rate was observed even for the small datasets with smaller dimensions. No difference was observed across different versions of the stray algorithm with different nearest neighbour searching mechanisms and their different implementations.

Figure 4 and Table 2 demonstrate how the stray algorithm outperforms the two versions of the HDoutliers algorithm under different circumstances. These limited set of examples were selected with the aim of highlighting some of the key feature of the stray algorithm:

- (1) All three algorithms were able to correctly capture the anomalous point at the rightmost upper corner of Figure 4 (a)- I, II, III). However, the second versions of the HDoutliers algorithm tend to generate some false positives, particularly with the small dimensions (Figure 4 (a)- II; Table 2, row 1 column 3)
- (2) Figure 4 (b)- III) shows its ability to deal with multimodal typical classes. The two clusters at the bottom of the graph represent two typical classes.

- Only the second version of the HDoutliers algorithm (Figure 4 (b)- II) that utilises the clustering step was able to detect the top-centred micro cluster that contains three anomalous data instances. However, forming small clusters prior to the distance calculation is not always helpful in detecting micro clusters.
- (3) Figure 4 (c)- II) (Table 2, row 3) shows a situation where even the second version of the HDoutliers algorithm fails in detecting micro clusters. The Leader algorithm in the HDoutliers algorithm uses a very small ball of a fixed radius to form clusters, and therefore, it now fails to capture the five points into a single cluster and instead generates three small clusters that are very close to one another. Both versions of the HDoutliers algorithm now fail to detect the micro cluster at the rightmost upper corner, because the dataset violates one of the major requirements of isolation of anomalous points or anomalous clusters. In stray, the value of k was set to 10. One can interpret the value of k as the maximum permissible size for a micro cluster. That is, for a small cluster to be a micro cluster, the number of data points in that cluster should be less than k . Otherwise, the cluster is considered a typical cluster.
- (4) Figure 4 (d)- III) demonstrates the ability of detecting inliers. The HDoutliers algorithm also has this ability of detecting inliers only when there are isolated inliers that are free from anomalous neighbours. Both versions of the HDoutliers algorithm fail to detect the two inliers since they are very close to one another and thereby jointly protect them as being anomalous (Table 2, row 4).
- (5) As explained in Section 3.2, Figure 4 (e)- II) shows how the clustering step of the second version of the HDoutliers algorithm can misguide the detection process and thereby increase the rate of false positives (Table 2, row 5 column 3). The dense areas of the dataset are marked with density curves. Two typical clusters are visible, one at the leftmost upper corner and the other at the rightmost bottom corner. An inlier is also

present in between the two typical classes. After forming cluster through the Leader algorithm, only one representative member is selected from each cluster for the nearest neighbour distance calculation. The selected member is now isolated and earns a very high anomalous score, leading the entire typical cluster at the rightmost bottom corner with 1,000 points to be identified as anomalous. In contrast, the stray algorithm is free from these problems because it does not involve any clustering step prior to the nearest neighbour distance calculation.

- (6) As explained in Section 3.2, Figure 4 (f)- II) shows how the clustering step can increase the rate of false negatives (Table 2, row 6 column 4). This dataset contains one typical class that is closely compacted in substance (the leftmost upper corner) and an obvious anomaly at the rightmost bottom corner. Since the typical class is a dense cluster, only a few data points are selected from the typical class for the nearest neighbour calculation. In this example, the clustering step substantially down-samples the original dataset, leading to a huge information loss in the representation of the original dataset. The blue dots in Figure 4 (f)- II) represent the selected members from each cluster for nearest neighbour calculations. Now, the reduced sample size is not enough for a proper calculation of the anomalous threshold based on extreme value theory.

6 Usage

We applied our stray algorithm to a dataset obtained from an automated pedestrian counting system with 43 sensors in the city of Melbourne, Australia ([City of Melbourne 2019](#), [Wang 2018](#)), to identify unusual pedestrian activities within the municipality. Identification of such unusual, critical behaviours of pedestrians at different city locations at different times of the day is important because it is a direct indication of a city's economic conditions, the related activities and the safety and convenience of the pedestrian experience ([City of Melbourne 2019](#)). It also guides and informs decision-making and planning. This

case study also illustrates how the stray algorithm can be used to deal with other data structures, such as temporal data and streaming data using feature engineering.

6.1 Handling Temporal Data

For this study, we consider the hourly pedestrian counts from January 2, 2019, to August 18, 2019. For clear visual illustration, Figure 5 shows only a limited part of the study period with the pedestrian counts at 43 locations in the city of Melbourne at different times of the day. The distribution of pedestrian counts follows a negatively skewed distribution. In general, pedestrian counts on weekdays display a bimodal distribution, while pedestrian counts on weekends follow a unimodal distribution. Now, the aim is to detect days with unusual behaviours. Since this involves a large collection of multivariate time series plots, each representing a day of the study period, manual monitoring is time-consuming and unusual behaviours are difficult to locate by visual inspection.

Detecting anomalous plots from a large collection of plots requires some pre-processing. In particular, to apply the stray algorithm, we need to convert this original dataset, with a large collection of multivariate time series plots, into a high dimensional dataset. A simple approach is to use features that describe the different shapes and patterns of the multivariate time series plots. Computing features that describe meaningful shapes and patterns in a given multivariate time series plot is straightforward with scagnostics (scatterplot diagnostics) developed by [Wilkinson et al. \(2005\)](#). For the current study, we select nine features: outlying (a measure of the proportion of long edges on all edges in the minimum spanning tree (MST)), skewed (a measure of skewness based on a ratio of quantiles of the edge lengths of MST), clumpy (a measure of clustering of data points based on MST), sparse (a measure to check whether points in a scatterplot are confined to a lattice or a small number of locations on the plane), striated (a measure of coherence of data points), convex (the ratio of the area of the alpha hull and the area of the convex hull), skinny (the ratio of perimeter to

area of a polygon), stringy (the ratio of diameter to length of a network) and monotonic (the squared Spearman correlation coefficient) [[Dang and Wilkinson \(2014\)](#); wilkinson2005graph]. Once we extract these nine features from each plot, we convert our original collection of multivariate time series plots into a dataset with nine dimensions and 228 data instances covering the study period from January 2, 2019, to August 18, 2019. Figure 6 shows the O3 plot (Overview of Outliers plot) ([Unwin 2019](#)) summarizing feature combinations on which those days are anomalies and on what groups of features have these days been identified as anomalies. There is a row for each feature combination for which anomalies are found. Two white columns separate the feature combinations and the anomalies detected. Each row of the block on the left shows which feature combination defines the row. There are 9 columns, one for each feature and a cell is gray if the feature is a part of the combination. From this analysis, 13 days were found to be anomalies in at least one of the sub feature spaces defined by different feature combinations. These anomalies are marked by red cells on the right block in Figure 6. Figure 7 provides feature-based representation of the original collection of multivariate time series plots. Each point in this high-dimensional data space corresponds to a single multivariate time series plot (or a day) in the original collection of multivariate time series plots. Anomalies determined by the stray algorithm in at least one of the sub feature spaces defined by different feature combinations are highlighted in Figure 7. The corresponding multivariate time series plots (or days) are highlighted in Figure 5. Visual inspection also confirms the anomalous behaviour of these individual multivariate time series plots. Most of these anomalous days display an unusual rise later in the day. Most of the anomalies in January (14, 15, 19 and 20 January 2019) cover the 2019 Australian Open, a Grand Slam tennis tournament that took place at Melbourne Park from 14 to 27 January 2019. This annual tennis tournament attracts many thousands of tennis fans from all around the worlds. Further investigations regarding 13 January 2019, reveal that there was a musical concert in Melbourne city and the unusual rise later in the day

could be due to the concert participants. Similar patterns were observed with the remaining anomalies detected.

After detecting the days with anomalous pedestrian behaviours, further investigation is carried out for each day to detect the locations with anomalous behaviours within the selected day. Once we focus on one day, we obtain a collection of 43 time series with hourly pedestrian counts generated from the 43 sensors located at different geographic locations in the city (Figure 8). For this analysis, we extract 11 time series features (similar to [Talagala et al. 2020](#), [Hyndman et al. 2015](#)) and convert the original collection of time series into a data space with 11 dimensions and 43 data instances (Figure 9). Now, each point in this high-dimensional space correspond to a single time series (or sensor) in Figure 8. The stray algorithm declares three points as anomalous points in this high dimensional data space. These points correspond to the sensors at Southbank, the Art centre and St Kilda Rd-Alexandra Gardens in Melbourne.

These types of findings play a critical role in making decisions about urban planning and management; to identify opportunities to improve city walkability and transport measures; to understand the impact of major events and other extreme conditions on pedestrian activity, and thereby assist in making decisions regarding security and resource requirements; and to plan and respond to emergency situations, etc.

6.2 Handling Streaming Data

Owing to the unsupervised nature of the stray algorithm, it can easily be extended for streaming data. A sliding window of fixed length can be used to deal with streaming data. Then, datasets in each window can be treated as a batch dataset ([Talagala, Hyndman, Leigh, Mengersen and Smith-Miles 2019](#)) and the stray algorithm can be applied to each window to detect anomalies in the datasets defined by the corresponding window.

It also can be used to identify anomalous time series within a large collection of streaming temporal data. Let $W[t, t + w]$ represent a sliding window containing n number of individual time series of length w . First, we extract m features (similar to [Hyndman et al. \(2015\)](#) and [Talagala, Hyndman, Leigh, Mengersen and Smith-Miles \(2019\)](#)) from each and every time series in this window. This step gives rise to an $n \times m$ feature matrix where each row now corresponds to a time series in the original collection of time series. Once we convert our original collection of time series into a high-dimensional dataset, we can apply the stray algorithm to identify anomalous points within this m -dimensional data space. The corresponding time series are then declared as anomalous series within the large collection of time series in the corresponding sliding window.

7 Conclusions and Further Research

The HDoutliers algorithm by [Wilkinson \(2017\)](#) is a powerful algorithm for detecting anomalies in high-dimensional data. However, it suffers from a few limitations that significantly hinder its ability to detect anomalies under certain situations. In this study, we propose an improved algorithm, the stray algorithm, that addresses these limitations. The stray algorithm has many special features: (1) It can deal with both one dimensional and high dimensional data as it is based on distance measures. By extracting k nearest neighbour distances for each data instance, it converts any high dimensional anomaly detection problem into a one dimensional problem. (2) Since the anomalous threshold calculation is a data driven approach, the algorithm is unsupervised in nature and therefore does not require labeled training datasets. (3) Most of the existing algorithms involve a manual anomalous threshold for binary classification as anomalies or typical points. Since the stray algorithm uses an anomalous threshold based on extreme value theory, it has a valid probabilistic interpretation. (4) It deals with masking problems and detects micro clusters as it does not limit its threshold calculation only to the nearest neighbour distances and instead uses k nearest neighbour distances. (5) Since it uses fast nearest neighbour searching

mechanisms it can be easily extended to streaming data using sliding windows. (6) Owing to the use of k nearest neighbour distances it can deal with multimodal distributions. (7) In addition to a label, the stray algorithm also assigns an anomalous score to each data instance to indicate the degree of outlieriness of each measurement. (8) Owing to the use of k nearest neighbour distances it also detect inliers, which is overlooked in most past research. We also demonstrate how the stray algorithm can assist in detecting anomalies present in other data structures using feature engineering.

While the HDoutliers algorithm is powerful, we have provided several classes of counterexamples in this paper where the structural properties of the data did not enable HDoutliers to detect certain types of outliers. We demonstrated on these counterexamples that the stray algorithm outperforms HDoutliers, in terms of both accuracy and computational time. It is certainly common practice to evaluate the strength of an algorithm using collections of test problems with various challenging properties. However, we acknowledge that these counterexamples are not diverse and challenging enough to enable us to comment about the unique strengths and weaknesses of these two algorithms, nor to generalise our findings to conclude that stray is always the superior algorithm. This study should be viewed as an attempt to simulate further investigation on the HDoutliers algorithm and its successors, with the ultimate goal to achieve further improvements across the entire problem space defined by various high-dimensional datasets. An important open research problem is therefore to assess the effectiveness of these algorithms across the the broadest possible problem space defined by different datasets with diverse properties (Kang et al. 2017). It is an interesting question to explore the impact of other classes of problems with various structural properties affect the performance of the stray algorithm and where its weaknesses might lie. This kind of instance space analysis (Smith-Miles et al. 2014) will enable further insights into improved algorithm design.

In our proposed algorithm, the anomalous threshold calculation was performed under the assumption that the distribution of k -nearest neighbours with the maximum gap is in the maximum domain of attraction of the Gumbel distribution. This requirement is satisfied by a wide range of distributions, ranging from those with light tails to moderately heavy tails that decrease to zero faster than any power function (Embrechts et al. 2013). Future studies will be required for alternative methods to find better values for the anomalous threshold in the presence of other sub-classes of EVT. For the current work, parameter k , which determines the size of the neighbourhood is introduced as a user-defined parameter that can be selected according to the application as too large values of k could increase the rate of false positives and too small values of k could increase the rate of false negatives. Further work will be needed to perform an optimization to select the best k as the performance of the algorithms can depend on the value of k and algorithms can reach their peak performance for different choices of k (Campos et al. 2016). One possibility is to adopt the method proposed by Zhang et al. (2017), for fast learning an optimal- k -value for each test sample. The features we used in this paper were directed by the data structures and typical patterns imposed by a given application. Domain specific knowledge plays a vital role when selecting suitable features for a given application.

In the simulation study, different versions of the algorithms (stray and HDoutlier) were applied on both real and simulated datasets. Some of the two-dimensional simulated data sets and real examples include high correlation among variables. But for very high dimensional data, each column is randomly generated from the standard normal distribution. One interesting issue is to look at the effect of the correlation on the algorithm performance when the number of dimensions of the datasets gets higher. But we leave this for future work as it is more complicated and not the main focus of this paper.

Anomaly detection problems commonly appear in many applications in different application domains. Therefore, it is hoped that different people with different knowledge levels will use the stray algorithm for many different purposes. Therefore, we expect future studies to develop interactive data visualisation tools that can enable exploring anomalies using a combination of graphical and numerical methods.

Supplementary Materials

Data and scripts:

Datasets and R code to reproduce all figures in this article (main.R and R package `stray` (Talagala, Hyndman and Smith-Miles 2019)).

R package `stray`:

The `stray` package consists of the implementation of the stray algorithm as described in this article. Version 0.1.1 of the package was used for the results presented in the article and is available from CRAN. The development version of the package is available from <https://github.com/pridiltal/stray>.

R-packages:

Each of the R packages used in this article (`ggplot2` (Wickham 2016), `dplyr` (Wickham et al. 2019), `tidyr` (Wickham and Henry 2019), `HDoutliers` (Fraley 2018), `lvplot` (Wickham and Hofmann 2016) are available online (URLs are provided in the bibliography).

Acknowledgements

This research was supported in part by the Monash eResearch Centre and eSolutions-Research Support Services through the use of the MonARCH (Monash Advanced Research Computing Hybrid) HPC Cluster. We thank Heike Hofmann, Eamonn Keogh and all the anonymous reviewers for their valuable comments and suggestions to improve our work. We also thank Anthony Unwin for his valuable feedback for our R implementation. Further, we thank Sevvandi

Kandanaarachchi and Mario A. Muñoz for joining the discussions during the initial stage of the project and Anastasios Panagiotelis for his valuable support during the revision of the manuscript. We also thank the Australian Centre of Excellence for Mathematical and Statistical Frontiers for supporting this work.

References

- Abuzaid, A., Hussin, A. and Mohamed, I. (2013), 'Detection of outliers in simple circular regression models using the mean circular error statistic', *Journal of Statistical Computation and Simulation* **83**(2), 269–277.
- Aggarwal, C. C. (2017), *Outlier analysis*, second edition. edn, Cham, Switzerland : Springer.
- Ben-Gal, I. (2005), Outlier detection, in 'Data mining and knowledge discovery handbook', Springer, pp. 131–146.
- Beygelzimer, A., Kakadet, S., Langford, J., Arya, S., Mount, D. and Li, S. (2019), *FNN: Fast Nearest Neighbor Search Algorithms and Applications*. R package version 1.1.3. URL: <https://CRAN.R-project.org/package=FNN>
- Breunig, M. M., Kriegel, H.-P., Ng, R. T. and Sander, J. (2000), Lof: identifying density-based local outliers, in 'ACM Sigmod Record', Vol. 29, ACM, pp. 93–104.
- Burridge, P. and Taylor, A. M. R. (2006), 'Additive outlier detection via extreme-value theory', *Journal of Time Series Analysis* **27**(5), 685–701.
- Campos, G. O., Zimek, A., Sander, J., Campello, R. J., Micenková, B., Schubert, E., Assent, I. and Houle, M. E. (2016), 'On the evaluation of unsupervised outlier detection: measures, datasets, and an empirical study', *Data Mining and Knowledge Discovery* **30**(4), 891–927.
- Chandola, V., Banerjee, A. and Kumar, V. (2009), 'Anomaly detection: A survey', *ACM Computing Surveys* **41**(3), 1–58.

City of Melbourne (2019), *Pedestrian Volume in Melbourne*. Last accessed 2019-07-23. URL: <http://www.pedestrian.melbourne.vic.gov.au>

Clifton, D. A., Huguency, S. and Tarassenko, L. (2011), 'Novelty detection with multivariate extreme value statistics', *Journal of Signal Processing Systems* **65**(3), 371–389.

Dang, T. N. and Wilkinson, L. (2014), 'Transforming scagnostics to reveal hidden features', *IEEE Transactions on Visualization and Computer Graphics* **20**(12), 1624–1632.

Elseberg, J., Magnenat, S., Siegwart, R. and Nüchter, A. (2012 *a*), 'Comparison of nearest-neighbor-search strategies and implementations for efficient shape registration', *Journal of Software Engineering for Robotics* **3**(1), 2–12.

Elseberg, J., Magnenat, S., Siegwart, R. and Nüchter, A. (2012 *b*), 'Comparison of nearest-neighbor-search strategies and implementations for efficient shape registration', *Journal of Software Engineering for Robotics (JOSER)* **3**(1), 2–12.

Embrechts, P., Klüppelberg, C. and Mikosch, T. (2013), *Modelling Extremal Events: for Insurance and Finance*, Stochastic Modelling and Applied Probability, Springer Berlin Heidelberg.

Fraley, C. (2018), *HDoutliers: Leland Wilkinson's Algorithm for Detecting Multidimensional Outliers*. R package version 1.0. URL: <https://CRAN.R-project.org/package=HDoutliers>

Galambos, J., Lechner, J. and Simiu, E. (2013), *Extreme Value Theory and Applications: Proceedings of the Conference on Extreme Value Theory and Applications, Volume 1 Gaithersburg Maryland 1993*, Extreme Value Theory and Applications: Proceedings of the Conference on Extreme Value Theory and Applications, Gaithersburg, Maryland, 1993, Springer US.

Gao, J., Hu, W., Zhang, Z. M., Zhang, X. and Wu, O. (2011), Rkof: robust kernel-based local outlier detection, *in* 'Pacific-Asia Conference on Knowledge Discovery and Data Mining', Springer, pp. 270–283.

Goldstein, M. and Uchida, S. (2016), 'A comparative evaluation of unsupervised anomaly detection algorithms for multivariate data', *PIOS ONE* **11**(4), e0152173.

Grubbs, F. E. (1969), 'Procedures for detecting outlying observations in samples', *Technometrics* **11**(1), 1–21.

Gupta, M., Gao, J., Aggarwal, C. C. and Han, J. (2014), 'Outlier detection for temporal data: A survey', *IEEE Transactions on Knowledge and Data Engineering* **26**(9), 2250–2267.

Hadi, A. S. (1992), 'Identifying multiple outliers in multivariate data', *Journal of the Royal Statistical Society: Series B (Methodological)* **54**(3), 761–771.

Hartigan, J. A. and Hartigan, J. (1975), *Clustering Algorithms*, Vol. 209, Wiley New York.

Hodge, V. and Austin, J. (2004), 'A survey of outlier detection methodologies', *Artificial Intelligence Review* **22**(2), 85–126.

Hofmann, H., Wickham, H. and Kafadar, K. (2017), 'Value plots: Boxplots for large data', *Journal of Computational and Graphical Statistics* **26**(3), 469–477.

Hyndman, R. J. (1996), 'Computing and graphing highest density regions', *The American Statistician* **50**(2), 120–126.

Hyndman, R. J., Wang, E. and Laptev, N. (2015), Large-scale unusual time series detection, *in* '2015 IEEE International Conference on Data Mining Workshop (ICDMW)', pp. 1616–1619.

Jin, W., Tung, A. K., Han, J. and Wang, W. (2006), Ranking outliers using symmetric neighborhood relationship, *in* 'Pacific-Asia Conference on Knowledge Discovery and Data Mining', Springer, pp. 577–593.

Kandanaarachchi, S., Munoz, M. A., Hyndman, R. J., Smith-Miles, K. et al. (2018), On normalization and algorithm selection for unsupervised outlier detection, Technical report, Monash University, Department of Econometrics and Business Statistics.

Kang, Y., Hyndman, R. J. and Smith-Miles, K. (2017), 'Visualising forecasting algorithm performance using time series instance spaces', *International Journal of Forecasting* **33**(2), 345–358.

Kanungo, T., Mount, D. M., Netanyahu, N. S., Piatko, C. D., Silverman, R. and Wu, A. Y. (2002), 'An efficient k-means clustering algorithm: Analysis and implementation', *IEEE Transactions on Pattern Analysis & Machine Intelligence* (7), 881–892.

Liu, S., Maljovec, D., Wang, B., Bremer, P.-T. and Pascucci, V. (2016), 'Visualizing high-dimensional data: Advances in the past decade', *IEEE Transactions on Visualization and Computer Graphics* **23**(3), 1249–1268.

Madsen, J. H. (2018), *DDoutlier: Distance and Density-Based Outlier Detection*. R package version 0.1.0. URL: <https://CRAN.R-project.org/package=DDoutlier>

Novotny, M. and Hauser, H. (2006), 'Outlier-preserving focus+ context visualization in parallel coordinates', *IEEE Transactions on Visualization and Computer Graphics* **12**(5), 893–900.

R Core Team (2019), *R: A Language and Environment for Statistical Computing*, R Foundation for Statistical Computing, Vienna, Austria. URL: <https://www.R-project.org/>

- Schwarz, K. T. (2008), *Wind dispersion of carbon dioxide leaking from underground sequestration, and outlier detection in eddy covariance data using extreme value theory*, ProQuest.
- Shahid, N., Naqvi, I. H. and Qaisar, S. B. (2015), 'Characteristics and classification of outlier detection techniques for wireless sensor networks in harsh environments: a survey', *Artificial Intelligence Review* **43**(2), 193–228.
- Smith-Miles, K., Baatar, D., Wreford, B. and Lewis, R. (2014), 'Towards objective measures of algorithm performance across instance space', *Computers & Operations Research* **45**, 12–24.
- Talagala, P. D., Hyndman, R. J., Leigh, C., Mengersen, K. and Smith-Miles, K. (2019), 'A feature-based procedure for detecting technical outliers in water-quality data from in situ sensors', *Water Resources Research* **55**(11), 8547–8568.
- Talagala, P. D., Hyndman, R. J. and Smith-Miles, K. (2019), *stray: Anomaly Detection in High Dimensional and Temporal Data*. R package version 0.1.1.
- Talagala, P. D., Hyndman, R. J., Smith-Miles, K., Kandanaarachchi, S. and Muñoz, M. A. (2020), 'Anomaly detection in streaming nonstationary temporal data', *Journal of Computational and Graphical Statistics* **29**(1), 13–27.
- Tang, J., Chen, Z., Fu, A. W.-C. and Cheung, D. W. (2002), Enhancing effectiveness of outlier detections for low density patterns, in 'Pacific-Asia Conference on Knowledge Discovery and Data Mining', Springer, pp. 535–548.
- Unwin, A. (2019), 'Multivariate outliers and the o3 plot', *Journal of Computational and Graphical Statistics* pp. 1–11.
- Wang, E. (2018), *rwalkr: API to Melbourne Pedestrian Data*. R package version 0.4.0. URL: <https://CRAN.R-project.org/package=rwalkr>

Weissman, I. (1978), 'Estimation of parameters and large quantiles based on the k largest observations', *Journal of the American Statistical Association* **73**(364), 812–815.

Wickham, H. (2016), *ggplot2: Elegant Graphics for Data Analysis*, Springer-Verlag New York. URL: <https://ggplot2.tidyverse.org>

Wickham, H., François, R., Henry, L. and Müller, K. (2019), *dplyr: A Grammar of Data Manipulation*. R package version 0.8.3. URL: <https://CRAN.R-project.org/package=dplyr>

Wickham, H. and Henry, L. (2019), *tidyr: Tidy Messy Data*. R package version 1.0.0. URL: <https://CRAN.R-project.org/package=tidyr>

Wickham, H. and Hofmann, H. (2016), *lvplot: Letter Value 'Boxplots'*. R package version 0.2.0. URL: <https://CRAN.R-project.org/package=lvplot>

Wilkinson, L. (2017), 'Visualizing big data outliers through distributed aggregation', *IEEE Transactions on Visualization and Computer Graphics* **24**(1), 256–266.

Wilkinson, L., Anand, A. and Grossman, R. (2005), Graph-theoretic scagnostics, in 'IEEE Symposium on Information Visualization, 2005. INFOVIS 2005.', IEEE, pp. 157–164.

Williams, K. T. (2016), Local parametric density-based outlier detection and ensemble learning with applications to malware detection, PhD thesis, The University of Texas at San Antonio.

Zhang, R. (2017), 'Performance of kd-tree vs brute-force nearest neighbor search on gpu?', Computational Science Stack Exchange.
URL: <https://scicomp.stackexchange.com/q/26873> (version: 2017-05-13).

Zhang, S., Li, X., Zong, M., Zhu, X. and Wang, R. (2017), 'Efficient knn classification with different numbers of nearest neighbors', *IEEE transactions on neural networks and learning systems* 29(5), 1774–1785.

Accepted Manuscript

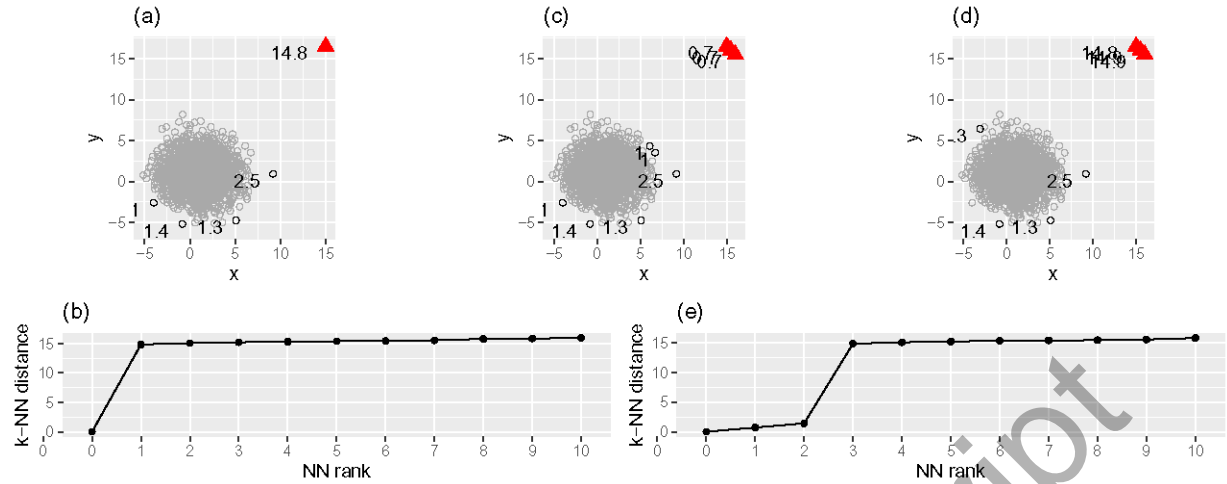


Fig. 1 Difference between the nearest neighbour distance and the k -nearest neighbour distance with the maximum gap. (a) Dataset contains only one anomaly at (15,16.5). Nearest neighbour distance are indicated by the labels in the figure. (b) Change in the k -nearest neighbour distances of the anomaly. (c) Dataset contains micro cluster around (15,16.5). Nearest neighbour distances are indicated by the labels in the figure. (d) Dataset contains micro cluster around (15,16.5). For the three anomalies, the third nearest neighbour distance (indicated by the labels in the figure) has the maximum gap. (e) Change in the k -nearest neighbour distances of an anomaly from micro cluster around (15,16.5). Anomalies are represented by triangles and the dots correspond to the typical behaviour. For this illustration two-dimensional datasets are selected to maximize the chances of obtaining insights via visualization.

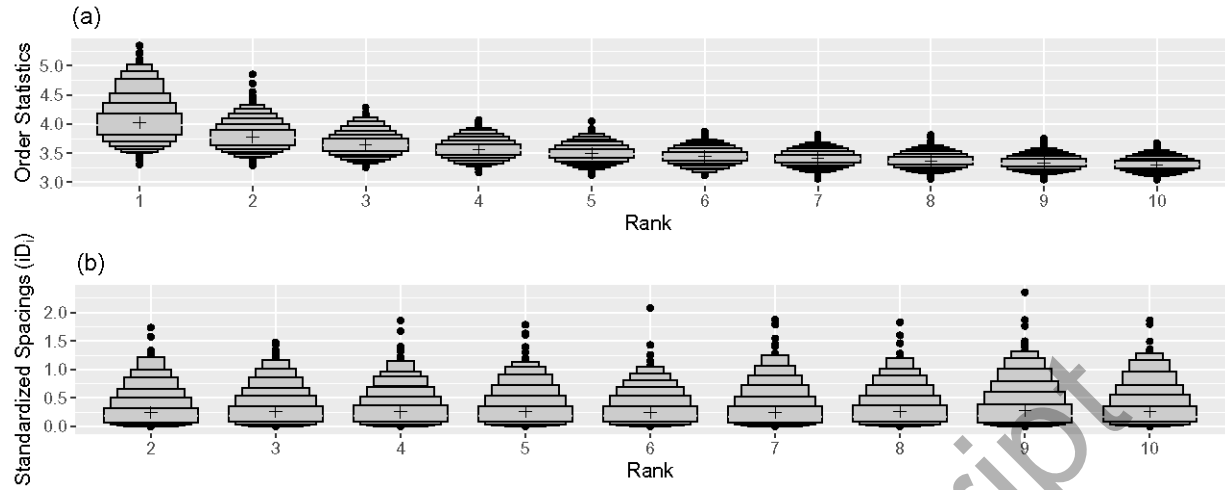


Fig. 2 (a) Distribution of the descending order statistics X_{in} and (b) distribution of the standardised spacings $iD_{i,n}$ for $i \in \{1, \dots, 10\}$ for 1, 000 samples each containing 20, 000 random numbers from the standard normal distribution.




Fig. 3 Scalability Performance. (a) HDoutliers algorithm without clustering step, (b-I) HDoutliers algorithm with clustering step, (c) stray algorithm with brute force nearest neighbour search using FNN R package implementation, (d) stray algorithm with kd-trees nearest neighbour search using 'FNN' R package implementation, (e) stray algorithm with brute force nearest neighbour search using 'nabor' R package implementation, (f) stray algorithm with kd-trees nearest neighbour search using 'nabor' R package implementation. For clear comparison, only a part of the measurements of the full experiment is displayed in (b-I). (b-II) presents the full version of (b-I). Black rectangle frame in (b-II) covers the plotting region of (b-I).

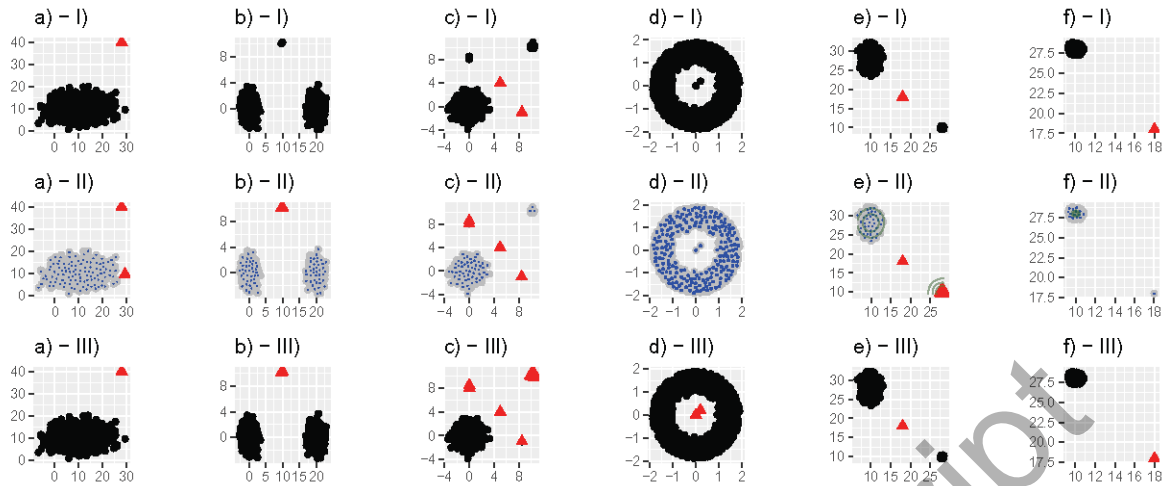


Fig. 4 Algorithm performance. (a) The top panel shows the results of the HDoutliers algorithm without a clustering step. (b) The middle panel shows the results of the HDoutliers algorithm with a clustering step. The representative member selected from each cluster formed by the Leader algorithm are marked as dark dots (c) The bottom panel shows the results of the improved algorithm with brute force k-nearest neighbour searching. The detected anomalies are marked as triangles. In each test, the critical value, α , was set to .1%. Two-dimensional datasets are selected to maximize the chances of obtaining insights via visualization.

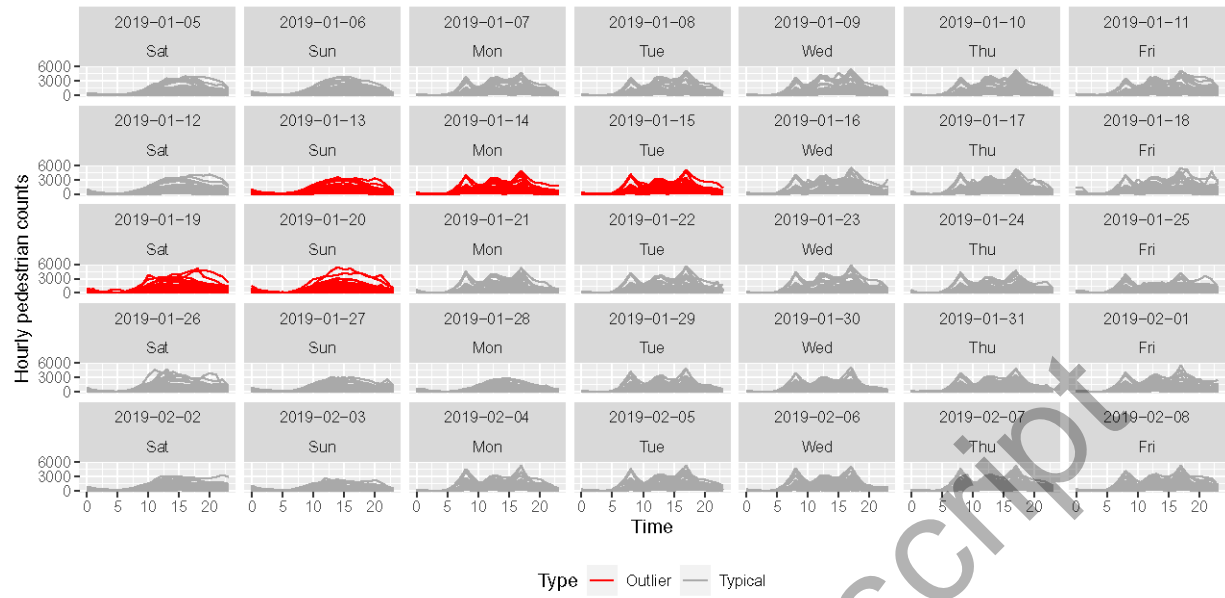


Fig. 5 Collection of multivariate time series plots of hourly pedestrian counts at 43 locations in the city Melbourne, Australia, from 2 January to 8 February 2019. Anomalous days detected by the stray algorithm using scagnostics are marked in dark color (red in the online version). This covers only a small part of the study period considered (from January 2, 2019, to August 18, 2019).

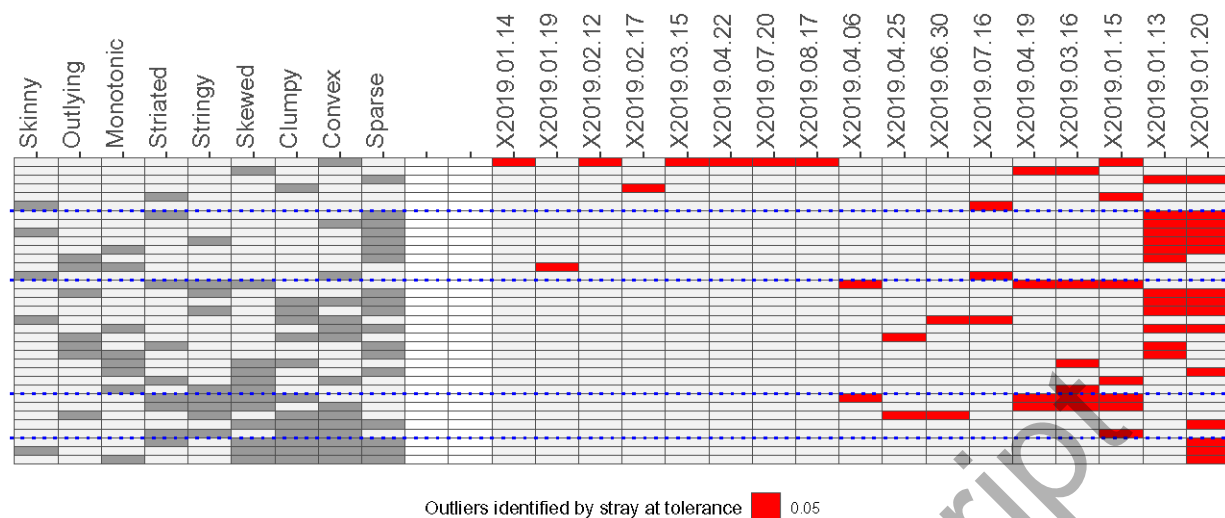


Fig. 6 O3 plot of data relating to hourly pedestrian counts at 43 locations in the city Melbourne, Australia, from January 2, 2019, to August 18, 2019. Thirteen days were found to be anomalies on some combination of features. Anomalous days detected by the stray algorithm are marked in dark cells (red in the online version). Two days were anomalies on several combinations, 13-01-2019 and 20-01-2019.




Fig. 7 Feature-based representation of the collection of multivariate time series plots using scagnostics. In each feature, anomalies determined by the stray algorithm in at least one of the sub feature spaces defined by different feature combinations are represented in dark colour (red in the online version). The columns of the data are normalised such that the data are bounded by the unit hypercube.




Fig. 8 Multivariate time series plot of hourly counts of pedestrians measured at 43 different sensors in the city of Melbourne, on 20 January 2019. The anomalous time series detected by the stray algorithm using time series features are marked in dashed lines).




Fig. 9 Feature-based representation of the collection of time series on 20 January 2019. In each plot, anomalies determined by the stray algorithm are represented in light (red in the online version) colour. The columns of the data are normalised such that the data are bounded by the unit hypercube.

Table 2 Performance metrics – False positive (FP) and False negative (FN) rates. HDo-WoC: HDoutliers algorithm without clustering step, HDo-WC: HDoutliers algorithm with clustering step.

| Data set | HDo-WoC FP | HDo-WoC FN | HDo-WC FP | HDo-WC FN | Stray FP | stray FN |
|----------|------------|------------|-----------|-----------|----------|----------|
| a | 0.00000 | 0.00000 | 0.00100 | 0.00000 | 0.00000 | 0.00000 |
| b | 0.00000 | 0.00150 | 0.00000 | 0.00000 | 0.00000 | 0.00000 |
| c | 0.00000 | 0.00694 | 0.00000 | 0.00496 | 0.00000 | 0.00000 |
| d | 0.00000 | 0.00200 | 0.00000 | 0.00200 | 0.00000 | 0.00000 |
| e | 0.00000 | 0.00000 | 0.49975 | 0.00000 | 0.00000 | 0.00000 |
| f | 0.00000 | 0.00000 | 0.00000 | 0.00050 | 0.00000 | 0.00000 |